

Project 1 - California Water Usage

Welcome to the first project in CS260! We will be exploring possible connections between water usage, geography, and income in California. The water data for this project was procured from the [California State Water Resources Control Board](#) and curated by the [Pacific Institute](#). The map data includes [US topography](#), [California counties](#), and [ZIP codes](#).

The dataset on income comes from the IRS ([documentation](#)). We have identified some interesting columns in the dataset, but a full description of all the columns (and a definition of the population in the dataset and some interesting anonymization procedures they used) is available in this [description](#).

As usual, **run the cell below** to prepare the automatic tests. **Passing the automatic tests does not guarantee full credit on any question.** The tests are provided to help catch some common errors, but it is *your* responsibility to answer the questions correctly.

In [1]:

```
# Run this cell, but please don't change it.

import numpy as np
import math
from datascience import *

# These lines set up the plotting functionality and formatting.
import matplotlib
# matplotlib.use('Agg', warn=False)
%matplotlib inline
import matplotlib.pyplot as plots
plots.style.use('fivethirtyeight')

# These lines load the tests.
from client.api.assignment import load_assignment
project1 = load_assignment('project1.ok')

=====
Assignment: Project 1
OK, version v1.18.1
=====
```

First, load the data. Loading may take some time.

In [2]:

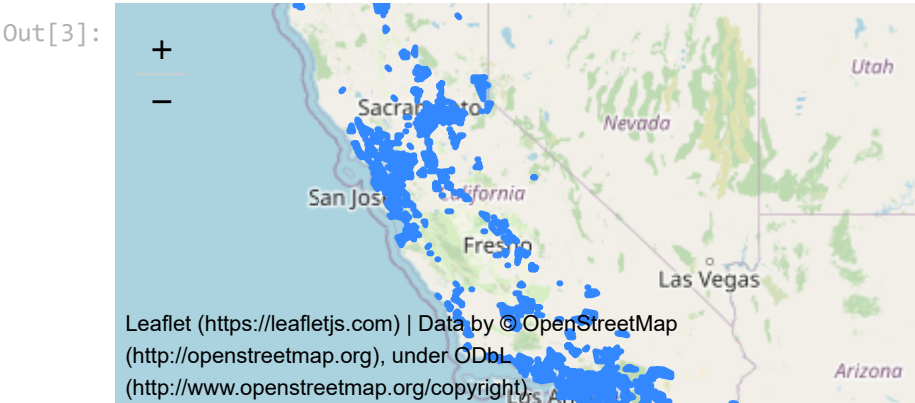
```
# Run this cell, but please don't change it.

districts = Map.read_geojson('water_districts.geojson')
zips = Map.read_geojson('ca_zips.geojson.gz')
usage_raw = Table.read_table('water_usage.csv', dtype={'pwsid': str})
income_raw = Table.read_table('ca_income_by_zip.csv', dtype={'ZIP': str}).drop(['STATEF
wd_vs_zip = Table.read_table('wd_vs_zip.csv', dtype={'PWSID': str, 'ZIP': str}).set_for
```

Part 0: Maps

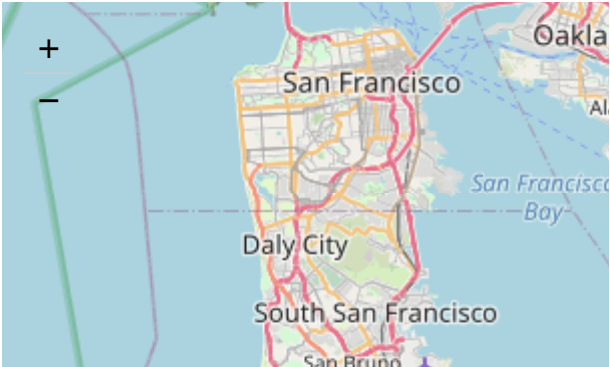
The `districts` and `zips` data sets are `Map` objects. Documentation on mapping in the `datascience` package can be found at data8.org/datascience/maps.html. To view a map of California's water districts, run the cell below. Click on a district to see its description.

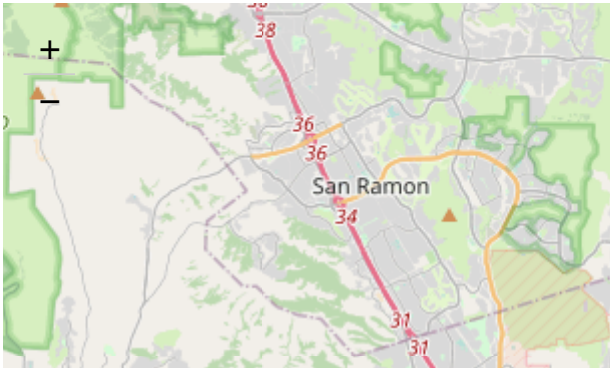

```
In [3]: districts.format(width=400, height=200)
```



A `Map` is a collection of regions and other features such as points and markers, each of which has a **string** `id` and various properties. You can view the features of the `districts` map as a table using `Table.from_records`.

```
In [4]: district_table = Table.from_records(districts.features)
district_table.show(3)
```

PWSID	feature id	popupContent
0110001		0 Alameda County Water District

PWSID	feature	id	popupContent
0110003		1	California Water Service Company Livermore
0110005		2	East Bay Municipal Utilities District

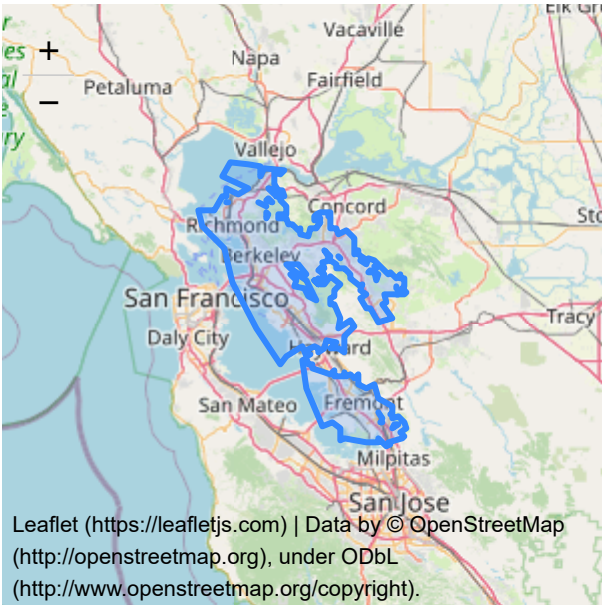
... (407 rows omitted)

To display a Map containing only two features from the `district_table` , call `Map` on a list containing those two features from the `feature` column.

Question 0.0 Draw a map of the Alameda County Water District (row 0) and the East Bay Municipal Utilities District (row 2).

```
In [5]: alameda_and_east_bay = district_table.take(0,2).column("feature")
Map(alameda_and_east_bay, height=300, width=300)
```

Out[5]:



Trust Notebook

Leaflet (<https://leafletjs.com>) | Data by © OpenStreetMap (<http://openstreetmap.org>), under ODbL (<http://www.openstreetmap.org/copyright>).

In [6]: `_ = project1.grade('q00')`

~~~~~  
Running tests

-----  
Test summary

Passed: 2

Failed: 0

[ooooooooook] 100.0% passed

*Hint:* If scrolling becomes slow on your computer, you can clear maps for the cells above by running  
Cell > All Output > Clear from the Cell menu.

## Part 1: California Income


Let's look at the `income_raw` table.

In [7]: `income_raw`

Out[7]:

| ZIP   | N1    | MARS1 | MARS2 | MARS4 | PREP  | N2    | NUMDEP | A00100 | N02650 | A02650 | N00200 |
|-------|-------|-------|-------|-------|-------|-------|--------|--------|--------|--------|--------|
| 90001 | 13100 | 6900  | 1890  | 4270  | 10740 | 29670 | 15200  | 181693 | 13100  | 184344 | 10220  |
| 90001 | 5900  | 1700  | 1970  | 2210  | 4960  | 17550 | 9690   | 203628 | 5900   | 204512 | 5610   |
| 90001 | 1480  | 330   | 760   | 390   | 1240  | 4710  | 2470   | 89065  | 1480   | 89344  | 1440   |
| 90001 | 330   | 50    | 210   | 70    | 290   | 1100  | 560    | 28395  | 330    | 28555  | 320    |
| 90001 | 160   | 30    | 100   | 40    | 130   | 510   | 250    | 24676  | 160    | 25017  | 150    |
| 90001 | 0     | 0     | 0     | 0     | 0     | 0     | 0      | 0      | 0      | 0      | 0      |
| 90002 | 12150 | 6330  | 1460  | 4330  | 9580  | 27240 | 14070  | 167261 | 12150  | 170095 | 9440   |
| 90002 | 5030  | 1510  | 1490  | 1980  | 4120  | 14410 | 7890   | 173280 | 5030   | 174335 | 4760   |
| 90002 | 1320  | 300   | 600   | 400   | 1060  | 4090  | 2180   | 78559  | 1320   | 78871  | 1270   |
| 90002 | 340   | 90    | 190   | 90    | 270   | 1060  | 530    | 28502  | 340    | 28558  | 320    |

... (8888 rows omitted)



Some observations:

1. The table contains several numerical columns and a column for the ZIP code.
2. For each ZIP code, there are 6 rows. Each row for a ZIP code has data from tax returns in one *income bracket* -- a group of people who make between some income and some other income.
3. According to the IRS documentation, all the numerical columns are *totals* -- either total numbers of returns that fall into various categories, or total amounts of money (in thousands of

dollars) from returns in those categories. For example, the column 'N02650' is the number of returns that included a total income amount, and 'A02650' is the total amount of total income (in thousands of dollars) from those returns.


**Question 1.0.** Since we don't care about income brackets, but we do care about totals per ZIP code, let's group together our income data by ZIP code. Assign the name `income_by_zipcode` to a table with just one row per ZIP code. When you group according to ZIP code, the remaining columns should be summed. In other words, for any other column such as 'N02650', the value of 'N02650' in a row corresponding to ZIP code 90210 (for example) should be the sum of the values of 'N02650' in the 6 rows of `income_raw` corresponding to ZIP code 90210.

```
In [8]: income_by_zipcode = income_raw.group("ZIP", sum)
income_by_zipcode
```

```
Out[8]:
```

| ZIP   | N1<br>sum | MARS1<br>sum | MARS2<br>sum | MARS4<br>sum | PREP<br>sum | N2<br>sum | NUMDEP<br>sum | A00100<br>sum | N02650<br>sum | A02650<br>sum |
|-------|-----------|--------------|--------------|--------------|-------------|-----------|---------------|---------------|---------------|---------------|
| 90001 | 20970     | 9010         | 4930         | 6980         | 17360       | 53540     | 28170         | 527457        | 20970         | 531772        |
| 90002 | 18960     | 8230         | 3830         | 6800         | 15120       | 47200     | 24850         | 462823        | 18960         | 467128        |
| 90003 | 26180     | 11310        | 5130         | 9640         | 20570       | 64470     | 33760         | 612733        | 26180         | 618848        |
| 90004 | 27360     | 15330        | 7000         | 4670         | 20260       | 51180     | 17800         | 1.61777e+06   | 27360         | 1.64943e+06   |
| 90005 | 15430     | 8550         | 3870         | 2830         | 11210       | 29910     | 11130         | 707020        | 15430         | 717290        |
| 90006 | 22630     | 11470        | 5400         | 5630         | 17840       | 47590     | 20210         | 563530        | 22630         | 571157        |
| 90007 | 11710     | 6350         | 2270         | 3020         | 8310        | 23380     | 9950          | 311779        | 11710         | 315581        |
| 90008 | 14710     | 8060         | 2310         | 4110         | 9990        | 27000     | 10310         | 662036        | 14710         | 668523        |
| 90010 | 2210      | 1270         | 690          | 210          | 1760        | 3790      | 960           | 314333        | 2210          | 320471        |
| 90011 | 36670     | 15540        | 8600         | 12390        | 30240       | 95640     | 51260         | 857731        | 36670         | 864961        |

... (1473 rows omitted)



```
In [9]: _ = project1.grade('q10')
```

```
~~~~~
Running tests
```

```

Test summary
```

```
Passed: 2
```

```
Failed: 0
```

```
[ooooooooo] 100.0% passed
```

**Question 1.1.** Relabel the columns in `income_by_zipcode` to match the labels in `income_raw`; you probably modified all the names slightly in the previous question.

*Hint:* Inspect `income_raw.labels` and `income_by_zipcode.labels` to find the differences you need to change.

*Hint 2:* Since there are many columns, it will be easier to relabel each of them by using a `for` statement. See [Section 3.2](#) of the textbook for details.

*Hint 3:* You can use the `replace` method of a string to remove excess content. See [lab02](#) for examples.


*Hint 4:* To create a new table from an existing table with one label replaced, use `relabelled`. To **change** a label in an existing table permanently, use `relabel`. Both methods take two arguments: the old label and the new label. You can solve this problem with either one, but `relabel` is simpler.

```
In [10]: for label in income_by_zipcode:
 income_by_zipcode.relabel(label, label.strip(" sum"))
 income_by_zipcode
```

```
Out[10]:
```

| ZIP   | N1    | MARS1 | MARS2 | MARS4 | PREP  | N2    | NUMDEP | A00100      | N02650 | A02650      | I |
|-------|-------|-------|-------|-------|-------|-------|--------|-------------|--------|-------------|---|
| 90001 | 20970 | 9010  | 4930  | 6980  | 17360 | 53540 | 28170  | 527457      | 20970  | 531772      |   |
| 90002 | 18960 | 8230  | 3830  | 6800  | 15120 | 47200 | 24850  | 462823      | 18960  | 467128      |   |
| 90003 | 26180 | 11310 | 5130  | 9640  | 20570 | 64470 | 33760  | 612733      | 26180  | 618848      |   |
| 90004 | 27360 | 15330 | 7000  | 4670  | 20260 | 51180 | 17800  | 1.61777e+06 | 27360  | 1.64943e+06 |   |
| 90005 | 15430 | 8550  | 3870  | 2830  | 11210 | 29910 | 11130  | 707020      | 15430  | 717290      |   |
| 90006 | 22630 | 11470 | 5400  | 5630  | 17840 | 47590 | 20210  | 563530      | 22630  | 571157      |   |
| 90007 | 11710 | 6350  | 2270  | 3020  | 8310  | 23380 | 9950   | 311779      | 11710  | 315581      |   |
| 90008 | 14710 | 8060  | 2310  | 4110  | 9990  | 27000 | 10310  | 662036      | 14710  | 668523      |   |
| 90010 | 2210  | 1270  | 690   | 210   | 1760  | 3790  | 960    | 314333      | 2210   | 320471      |   |
| 90011 | 36670 | 15540 | 8600  | 12390 | 30240 | 95640 | 51260  | 857731      | 36670  | 864961      |   |

... (1473 rows omitted)



```
In [11]: _ = project1.grade('q11')
```

```
~~~~~
Running tests

-----
Test summary
  Passed: 1
  Failed: 0
[ooooooooook] 100.0% passed
```

**Question 1.2.** Create a table called `income` with one row per ZIP code and the following columns.

1. A `ZIP` column with the same contents as `'ZIP'` from `income_by_zipcode`.
2. A `returns` column containing the total number of tax returns that include a total income amount (column `'N02650'` from `income_by_zipcode`).
3. A `total` column containing the total income in all tax returns in thousands of dollars (column `'A02650'` from `income_by_zipcode`).
4. A `farmers` column containing the number of farmer returns (column `'SCHF'` from `income_by_zipcode`).

```
In [12]: income = Table().with_columns([
    'ZIP', income_by_zipcode.column('ZIP'),
    'returns', income_by_zipcode.column('N02650'),
    'total', income_by_zipcode.column('A02650'),
    'farmers', income_by_zipcode.column('SCHF')
])
income.set_format('total', NumberFormatter(0)).show(5)
```

| ZIP   | returns | total     | farmers |
|-------|---------|-----------|---------|
| 90001 | 20970   | 531,772   | 0       |
| 90002 | 18960   | 467,128   | 0       |
| 90003 | 26180   | 618,848   | 0       |
| 90004 | 27360   | 1,649,431 | 0       |
| 90005 | 15430   | 717,290   | 0       |

... (1478 rows omitted)

```
In [13]: _ = project1.grade('q12')
```

~~~~~  
Running tests

Question > Suite 1 > Case 2

```
>>> print(income.take([0, 1, 2]))
ZIP | returns | total | farmers
90001 | 20970 | 531,772 | 0
90002 | 18960 | 467,128 | 0
90003 | 26180 | 618,848 | 0
```

Error: expected

```
# ZIP | returns | total | farmers
# 90001 | 20970 | 531,772 | 0
# 90002 | 18960 | 467,128 | 0
# 90003 | 26180 | 618,848 | 0
```

but got

```
# ZIP | returns | total | farmers
# 90001 | 20970 | 531,772 | 0
# 90002 | 18960 | 467,128 | 0
# 90003 | 26180 | 618,848 | 0
```

Run only this test case with "python3 ok -q q12 --suite 1 --case 2"

```
Test summary
  Passed: 1
  Failed: 1
[ooooooooo] 50.0% passed
```

Question 1.3. What is the average total income reported on all California tax returns that include a total income amount? **Express the answer in *dollars* as an `int` rounded to the nearest dollar.**

```
In [14]: average_income = int(np.average(income.column("total")))
         average_income
```

```
Out[14]: 828283
```

Question 1.4. All ZIP codes with less than 100 returns (or some other special conditions) are grouped together into one ZIP code with a special code. Remove the row for that ZIP code from the `income` table. *Hint:* This ZIP code value has far more returns than any of the other ZIP codes.

Hint: To **remove** a row in the `income` table using `where`, assign `income` to the smaller table using the following expression structure:

```
income = income.where(...)
```

Hint 2: Each ZIP code is represented as a string, not an `int`.

```
In [15]: income = income.where("ZIP", are.not_equal_to('99999'))
```

```
In [16]: _ = project1.grade('q14')
```

```
~~~~~
Running tests
```

```
-----
Test summary
  Passed: 2
  Failed: 0
[oooooooooooo] 100.0% passed
```

Question 1.5. Among the tax returns in California for ZIP codes represented in the `incomes` table, is there an association between income and living in a ZIP code with a higher-than-average proportion of farmers?

Answer the question by comparing the average incomes for two groups of *tax returns*: those in ZIP codes with a greater-than-average proportion of farmers and those in ZIP codes with a less-than-average (or average) proportion. Make sure both of these values are displayed (preferably in a table). *Then, describe your findings.*

```
In [17]: # Build and display a table with two rows:
         # 1) incomes of returns in ZIP codes with a greater-than-average proportion of farmer
         # 2) incomes of returns in other ZIP codes
```



```

farm_proportion = income.column('farmers')/income.column('returns')
avg_farm_proportion = np.average(farm_proportion)
fincome = income.with_columns('farm_proportion', farm_proportion)
fincome1= fincome.where("farm_proportion", are.above(avg_farm_proportion))
fincome_avg1 = (np.average(fincome1.column('total')/fincome1.column('returns'))*1000)
fincome2= fincome.where('farm_proportion', are.below_or_equal_to(avg_farm_proportion))
fincome_avg2 = np.average((fincome2.column('total')/fincome2.column('returns'))*1000)
income_by_farm_proportion = Table().with_columns('Averages', make_array(fincome_avg1,
                                                                           fincome_avg2))

income_by_farm_proportion

```

Out[17]: **Averages**

60225.7

79038.3

There is a 13.5% disparity between incomes of returns in ZIP codes with a greater-than-average proportion of farmers vs. incomes of returns in other ZIP codes. This can be contributed to other skill professionals that are included in the greater-than-average proportion of farmers.

Question 1.6. Investigate the same question by comparing two histograms: the average incomes of ZIP codes that have above-average vs below-average proportions of farmers. Quantify and describe the difference in the [standard deviations](#) of average incomes for the two kinds of ZIP codes.

In [18]:


```

# You do not need to change this cell; just look at the chart it generates.

bins = np.arange(20000, 300000, 5000)
avg_income = 1000 * income.column('total')/income.column('returns')
binned_incomes=Table()\
    .with_column('income', avg_income)\
    .where(farm_proportion > np.mean(farm_proportion))\
    .bin(bins=bins).relabeled(1, 'Avg. income in ZIP codes with above-average farmer pr
    .join('bin', Table().with_column('income', avg_income)\
    .where(farm_proportion <= np.mean(farm_proportion))\
    .bin(bins=bins).relabeled(1, 'Avg. income of ZIP codes with below-average farmer
binned_incomes.hist(counts='bin', bins=bins, unit='dollar')

```

C:\Anaconda\lib\site-packages\datascience\tables.py:4732: UserWarning: counts arg of hist is deprecated; use bin_column
 warnings.warn("counts arg of hist is deprecated; use bin_column")
C:\Anaconda\lib\site-packages\datascience\tables.py:5206: UserWarning: FixedFormatter should only be used together with FixedLocator
 axis.set_xticklabels(ticks, rotation='vertical')



```
In [19]: # Compute and compare the spread of both distributions
SD_1 = np.std(binned_incomes.column(
    "Avg. income in ZIP codes with above-average farmer proportion"))
SD_2 = np.std(binned_incomes.column(
    "Avg. income of ZIP codes with below-average farmer proportion"))
Standard_Deviation = Table().with_columns("Standard_Deviation", make_array(SD_1, SD_2))
Standard_Deviation
```


Out[19]: **Standard_Deviation**

9.88413

30.3214

The standard deviation (SD) for above and below average farmer proportion follows a similar theme in the explanation for the disparity between incomes. The bell shape curve/distribution for above average farmer proportion is narrower (less of a spread) vs. the below average farmer proportion, where the spread is wider. We can see these points by looking at the horizontal axis above. In addition, the population also should be factored into the explanation. We can also see from the zip codes coverage that certain regions in California are heavily populated (e.g., LA and the Bay). There are higher income earners in these areas but with the increase population contributes to the spread for above average. However, the remaining regions in California are less populated with potentially more farmers with lower income.

ZIP codes cover all the land in California and do not overlap. Here's a map of all of them.



Question 1.7. Among the ZIP codes represented in the `incomes` table, is there an association between high average income and some aspect of the ZIP code's location? If so, describe one aspect of the location that is clearly associated with high income.

Answer the question by drawing a map of all ZIP codes that have an average income above 100,000 dollars. *Then, describe an association that you observe.*

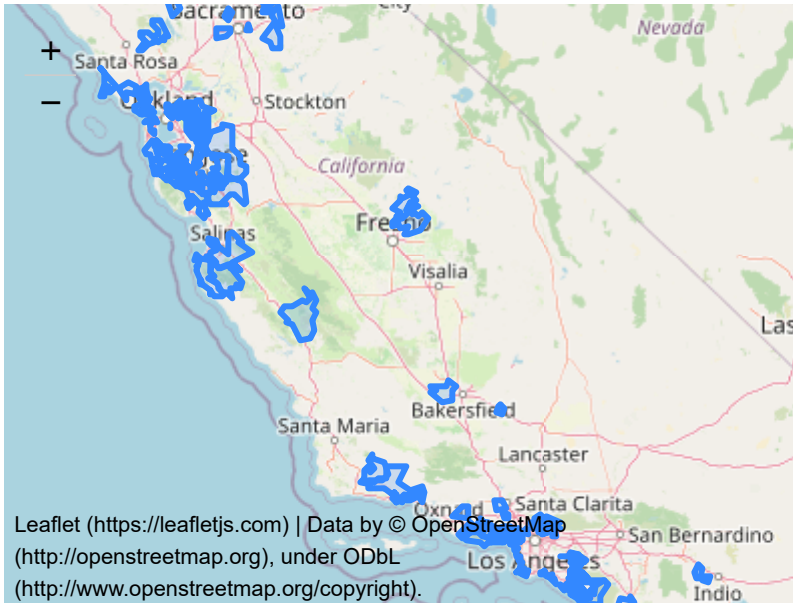
In order to create a map of certain ZIP codes, you need to

- Construct a table containing only the ZIP codes of interest, called `high_average_zips`,
- Join `high_average_zips` with the `zip_features` table to find the region for each ZIP code of interest,
- Call `Map(...)` on the column of features (provided).

In [20]:

```
# Write code to draw a map of only the high-income ZIP codes
zip_features = Table.from_records(zips.features)
high_average_zips1 = income.with_columns('avg', ((income.column(
    "total")/income.column("returns"))*1000))
high_average_zips = high_average_zips1.where('avg', are.above(100000))
high_zips_with_region = high_average_zips.join('ZIP', zip_features, "ZIP")
Map(list(high_zips_with_region.column('feature')), width=400, height=300)
```

Out[20]:



The higher average income locations are on the cost line of California and centralized with LA and the Bay area.

In [21]:

```
_ = project1.grade('q17')
```

~~~~~  
Running tests

-----  
Test summary

Passed: 2

Failed: 0

[ooooooooo] 100.0% passed

## Part 2: Water Usage

We will now investigate water usage in California. The `usage` table contains three columns:

- PWSID : The Public Water Supply Identifier of the district
- Population : Estimate of average population served in 2015
- Water : Average residential water use (gallons per person per day) in 2014-2015

In [22]:

```
# Run this cell to create the usage table
```

```
usage_raw.set_format(4, NumberFormatter)
max_pop = usage_raw.select([0, 'population']).group(0, max).relabel(1, 'Population')
avg_water = usage_raw.select([0, 'res_gpcd']).group(0, np.mean).relabel(1, 'Water')
usage = max_pop.join('pwsid', avg_water).relabel(0, 'PWSID')
usage
```

Out[22]:

| PWSID   | Population | Water   |
|---------|------------|---------|
| 0110001 | 340000     | 70.7    |
| 0110003 | 57450      | 90.2727 |
| 0110005 | 1390000    | 76      |
| 0110006 | 151037     | 57.1818 |
| 0110008 | 73067      | 96.6364 |
| 0110009 | 79547      | 68.6364 |
| 0110011 | 31994      | 85.8182 |
| 0310003 | 23347      | 82.8182 |
| 0410002 | 101447     | 142     |
| 0410005 | 11208      | 88.8182 |

... (401 rows omitted)

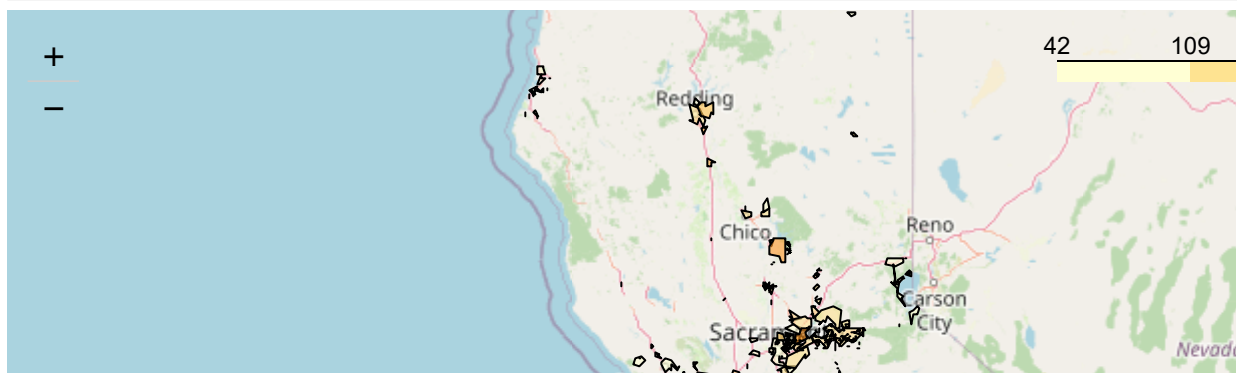
**Question 2.1.** Draw a map of the water districts, colored by the per capita water usage in each district.

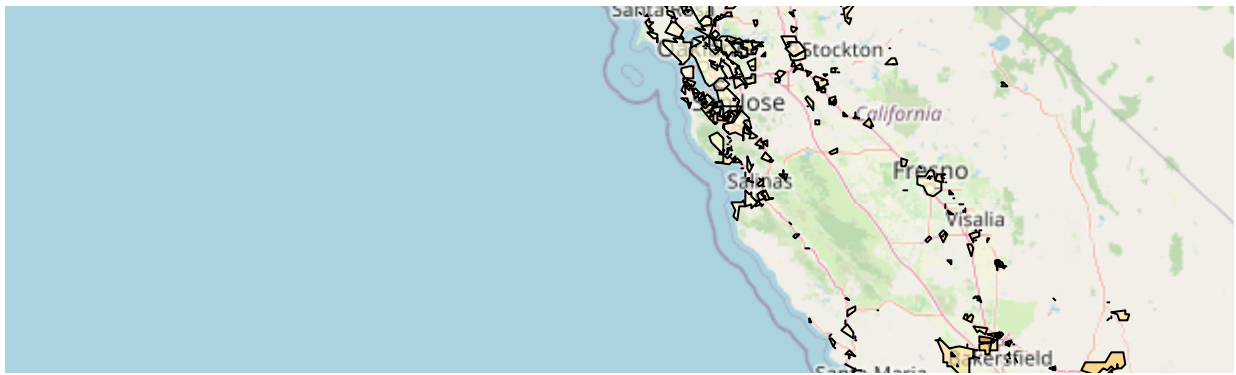
Use the `districts.color(...)` method to generate the map. It takes as its first argument a two-column table with one row per district that has the district PWSID as its first column. The label of the second column is used in the legend of the map, and the values are used to color each region.

In [23]:

```
per_capita_usage = Table().with_columns("PWSID", usage.column(0), "Per cap usage",
                                         usage.column(1*2))
districts.color(per_capita_usage, key_on='feature.properties.PWSID')
```

Out[23]:





```
In [24]: _ = project1.grade('q21')
```

```
~~~~~
Running tests
```

```

Test summary
```

```
 Passed: 2
```

```
 Failed: 0
```

```
[ooooooooook] 100.0% passed
```

**Question 2.2.** Based on the map above, which part of California appears to use more water per person, the San Francisco area or the Los Angeles area?

*Replace this line with a description of your findings.*

Next, we will try to match each ZIP code with a water district. ZIP code boundaries do not always line up with water districts, and one water district often covers multiple ZIP codes, so this process is imprecise. It is even the case that some water districts overlap each other. Nonetheless, we can continue our analysis by matching each ZIP code to the water district with the largest geographic overlap.

The table `wd_vs_zip` describes the proportion of land in each ZIP code that is contained in each water district and vis versa. (The proportions are approximate because they do not correctly account for discontinuous districts, but they're mostly accurate.)

```
In [25]: wd_vs_zip.show(10)
```

| PWSID   | ZIP   | District in ZIP | ZIP in District |
|---------|-------|-----------------|-----------------|
| 0110001 | 94536 | 9.41%           | 68.51%          |
| 0110001 | 94538 | 18.87%          | 67.31%          |
| 0110001 | 94539 | 13.13%          | 44.36%          |
| 0110005 | 94541 | 1.61%           | 68.11%          |
| 0110006 | 94541 | 18.68%          | 98.46%          |
| 0110005 | 94542 | 0.17%           | 6.79%           |
| 0110006 | 94542 | 18.24%          | 91.80%          |
| 0110001 | 94544 | 0.33%           | 3.02%           |

| PWSID   | ZIP   | District in ZIP | ZIP in District |
|---------|-------|-----------------|-----------------|
| 0110005 | 94544 | 0.00%           | 0.01%           |
| 0110006 | 94544 | 28.24%          | 97.61%          |

... (2836 rows omitted)

**Question 2.3.** Complete the `district_for_zip` function that takes a ZIP code. It returns the PWSID with the largest value of `ZIP in District` for that `zip_code`, if that value is at least 50%. Otherwise, it returns the string `'No District'`.

```
In [26]: def district_for_zip(zip_code):
zip_code = str(zip_code) # Ensure that the ZIP code is a string, not an integer
districts = wd_vs_zip.where('ZIP', are.equal_to(zip_code))
at_least_half = np.any(districts.column("ZIP in District")>=.5)
if at_least_half:
 return districts.sort("ZIP in District", descending=True).column(
 "PWSID").item(0)
else:
 return 'No District'

district_for_zip(94709)
```

Out[26]: '0110005'

```
In [27]: _ = project1.grade('q23')
```

~~~~~  
Running tests

-----  
Test summary

Passed: 4

Failed: 0

[ooooooooook] 100.0% passed

This function can be used to associate each ZIP code in the `income` table with a PWSID and discard ZIP codes that do not lie (mostly) in a water district.

```
In [28]: zip_pwsids = income.apply(district_for_zip, 'ZIP')
income_with_pwsid = income.with_column('PWSID', zip_pwsids).where(zip_pwsids != "No Dis
income_with_pwsid.set_format(2, NumberFormatter(0)).show(5)
```

| ZIP   | returns | total     | farmers | PWSID   |
|-------|---------|-----------|---------|---------|
| 90001 | 20970   | 531,772   | 0       | 1910067 |
| 90022 | 26680   | 767,484   | 0       | 1910036 |
| 90024 | 14690   | 4,395,487 | 20      | 1910067 |
| 90025 | 25110   | 4,019,082 | 20      | 1910067 |
| 90034 | 29950   | 1,828,572 | 0       | 1910067 |

... (662 rows omitted)

**Question 2.4.** Create a table called `district_data` with one row per PWSID and the following columns:

- **PWSID** : The ID of the district
- **Population** : Population estimate
- **Water** : Total annual water usage (millions of gallons)
- **Income** : Average income in dollars of all tax returns in ZIP codes that are (mostly) contained in the district according to `income_with_pwsid`.

*Hint:* First create a `district_income` table that sums the incomes and returns for ZIP codes in each water district.

```
In [29]: district_income = income_with_pwsid.group("PWSID", sum).join(
 'PWSID', usage, 'PWSID').drop ("ZIP sum", "farmers sum")
district_data = district_income.with_columns("Income", ((district_income.column(
 "total sum") / district_income.column("returns sum"))*1000)).drop ("total sum", "re
district_data.set_format(['Population', 'Water', 'Income'], NumberFormatter(0))
```

```
Out[29]:
```

| PWSID   | Population | Water | Income  |
|---------|------------|-------|---------|
| 0110001 | 340,000    | 71    | 79,032  |
| 0110005 | 1,390,000  | 76    | 82,497  |
| 0110006 | 151,037    | 57    | 52,924  |
| 0110008 | 73,067     | 97    | 163,257 |
| 0110009 | 79,547     | 69    | 133,902 |
| 0410002 | 101,447    | 142   | 50,401  |
| 0410006 | 18,300     | 286   | 38,721  |
| 0410011 | 9,615      | 92    | 44,707  |
| 0710001 | 106,455    | 110   | 53,551  |
| 0710003 | 197,536    | 102   | 73,914  |

... (200 rows omitted)

```
In [30]: _ = project1.grade('q24')
```

~~~~~  
Running tests

-----  
Test summary

Passed: 2

Failed: 0

[ooooooooook] 100.0% passed

**Question 2.5.** The `bay_districts` table gives the names of all water districts in the San Francisco

Bay Area. Is there an association between water usage and income among Bay Area water districts? Use the tables you have created to compare water usage between the 10 Bay Area water districts with the highest average income and the rest, then describe the association. *Do not include any districts in your analysis for which you do not have income information.*

The names below are just suggestions; you may perform the analysis in any way you wish.

```
In [31]: bay_districts = Table.read_table('bay_districts.csv')
bay_water_vs_income = bay_districts.join(
 "District", district_table, "popupContent").drop("feature").drop('id').join(
 "PWSID", district_data, "PWSID")
top_10 = bay_water_vs_income.sort("Income", descending=True).take(np.arange(10))
Bottom_19 = bay_water_vs_income.sort("Income").take(np.arange(19))
Gallons = np.average(top_10.column("Water"))-np.average(Bottom_19.column("Water"))
Gallons
```

Out[31]: 24.64545454545454

*Complete this one-sentence conclusion:* In the Bay Area, people in the top 10 highest-income water districts used an average of 24.6 more gallons of water per person per day than people in the rest of the districts.

```
In [32]: bay_districts = Table.read_table('bay_districts.csv')
bay_districts.join("District", district_table, "popupContent").drop (
 "feature", "id").join('PWSID', district_data, 'PWSID').sort (
 "Income", descending=True)
```

Out[32]:

| PWSID   | District                                            | Population | Water | Income    |
|---------|-----------------------------------------------------|------------|-------|-----------|
| 4110006 | California Water Service Company Bear Gulch         | 58,895     | 170   | 1,160,494 |
| 4310001 | California Water Service Company Los Altos/Suburban | 68,163     | 125   | 349,183   |
| 4310009 | Palo Alto, City of                                  | 64,403     | 89    | 334,057   |
| 4110017 | Menlo Park, City of                                 | 16,066     | 75    | 278,733   |
| 2110002 | Marin Municipal Water District                      | 188,200    | 86    | 176,643   |
| 0110008 | Pleasanton, City of                                 | 73,067     | 97    | 163,257   |
| 4310007 | Mountain View, City of                              | 76,781     | 66    | 138,570   |
| 4110021 | Estero Municipal Improvement District               | 37,165     | 66    | 137,893   |
| 4110001 | Mid-Peninsula Water District                        | 26,730     | 84    | 137,537   |
| 4110008 | California Water Service Company Mid Peninsula      | 135,918    | 70    | 135,967   |

... (19 rows omitted)

**Question 2.6.** In one paragraph, summarize what you have discovered through the analyses in this project and suggest what analysis should be conducted next to better understand California water usage, income, and geography. What additional data would be helpful in performing this next analysis?

The summary/analysis for this project would be a combination from the previous explanations. The



major item to call out would be regions in California that are heavily populated clearly sway favorable to income but not so favorable to water usage. Since there is an increase of ~24.6 gallons per person per day vs. the rest of the districts.

I think it would be helpful to know within populated zip codes (e.g. LA and the Bay) working class vs. government aid individuals. This is to potentially exclude individuals that receives government aid to know if that would change water usage per person which I suspect is a diluted number. I think this would exclude a good portion of the number of people who file a return in these zip codes. Which in return would increase the water usage.

Congratulations - you've finished Project 1 of CS260!

To submit:

1. Select `Run All` from the `Cell` menu to ensure that you have executed all cells, including the test cells. Make sure that the visualizations you create are actually displayed.
2. Select `Download as PDF via LaTeX (.pdf)` from the `File` menu. (Sometimes that seems to fail. If it does, you can download as HTML, open the .html file in your browser, and print it to a PDF.)
3. Read that file! If any of your lines are too long and get cut off, we won't be able to see them, so break them up into multiple lines and download again. If maps do not appear in the output, that's ok.
4. Submit that downloaded file (called `project1.pdf`) to Gradescope.

In [33]:

```
For your convenience, you can run this cell to run all the tests at once!
import os
_ = [project1.grade(q[:-3]) for q in os.listdir("tests") if q.startswith('q')]
```

```
~~~~~
Running tests
```

```
-----
Test summary
  Passed: 2
  Failed: 0
[ooooooooook] 100.0% passed
```

```
~~~~~
Running tests
```

```

Test summary
 Passed: 2
 Failed: 0
[ooooooooook] 100.0% passed
```

```
~~~~~
Running tests
```

```
-----
Test summary
  Passed: 1
  Failed: 0
[ooooooooook] 100.0% passed
```

~~~~~

Running tests

Question > Suite 1 > Case 2

```
>>> print(income.take([0, 1, 2]))
ZIP | returns | total | farmers
90001 | 20970 | 531,772 | 0
90002 | 18960 | 467,128 | 0
90003 | 26180 | 618,848 | 0
```

```
# Error: expected
#   ZIP | returns | total | farmers
#   90001 | 20970 | 531,772 | 0
#   90002 | 18960 | 467,128 | 0
#   90003 | 26180 | 618,848 | 0
# but got
#   ZIP | returns | total | farmers
#   90001 | 20970 | 531,772 | 0
#   90002 | 18960 | 467,128 | 0
#   90003 | 26180 | 618,848 | 0
```

Run only this test case with "python3 ok -q q12 --suite 1 --case 2"

Test summary
 Passed: 1
 Failed: 1
[ooooook.....] 50.0% passed

~~~~~

Running tests

-----  
Test summary  
    Passed: 2  
    Failed: 0  
[ooooooooook] 100.0% passed

~~~~~

Running tests

Test summary
 Passed: 2
 Failed: 0
[ooooooooook] 100.0% passed

~~~~~

Running tests

-----  
Test summary  
    Passed: 2  
    Failed: 0  
[ooooooooook] 100.0% passed

~~~~~

Running tests

```
-----  
Test summary  
  Passed: 4  
  Failed: 0  
[ooooooooook] 100.0% passed
```

```
~~~~~  
Running tests
```

```
-----  
Test summary  
  Passed: 2  
  Failed: 0  
[ooooooooook] 100.0% passed
```

If you want, draw some more maps below.

```
In [68]: # Your extensions here (completely optional)
```